

# Assignment 1

Aryaman Srivastava

2210110206

## Implementing Vector Clocks

### 1. Introduction

This report presents a comparative analysis of two vector clock implementations in distributed systems:

**Standard Vector Clock (VC)** - The traditional approach where complete vector clocks are transmitted with every message

**Singhal-Kshemkalyani (SK) Optimization** - An optimized approach that transmits only relevant vector clock entries

Our experiments demonstrate that the SK optimization achieves significant reduction in message overhead, with the optimization transmitting approximately 67-74% fewer entries per message compared to the standard approach.

### 2. Parameters

**Number of processes (n):** Varied from 10 to 15 (increments of 1)

**Inter-event time ( $\lambda$ ):** 5 ms

**Event ratio ( $\alpha$ ):** 1.5 (internal to message send events)

**Messages per process (m):** 50

**Topology:** Fully connected topology

In this topology:

- Every process is connected to every other process
- Each process  $i$  has edges to all processes  $j$  where  $j \neq i$
- For  $n$  processes, each process has  $(n-1)$  neighbors

For example:

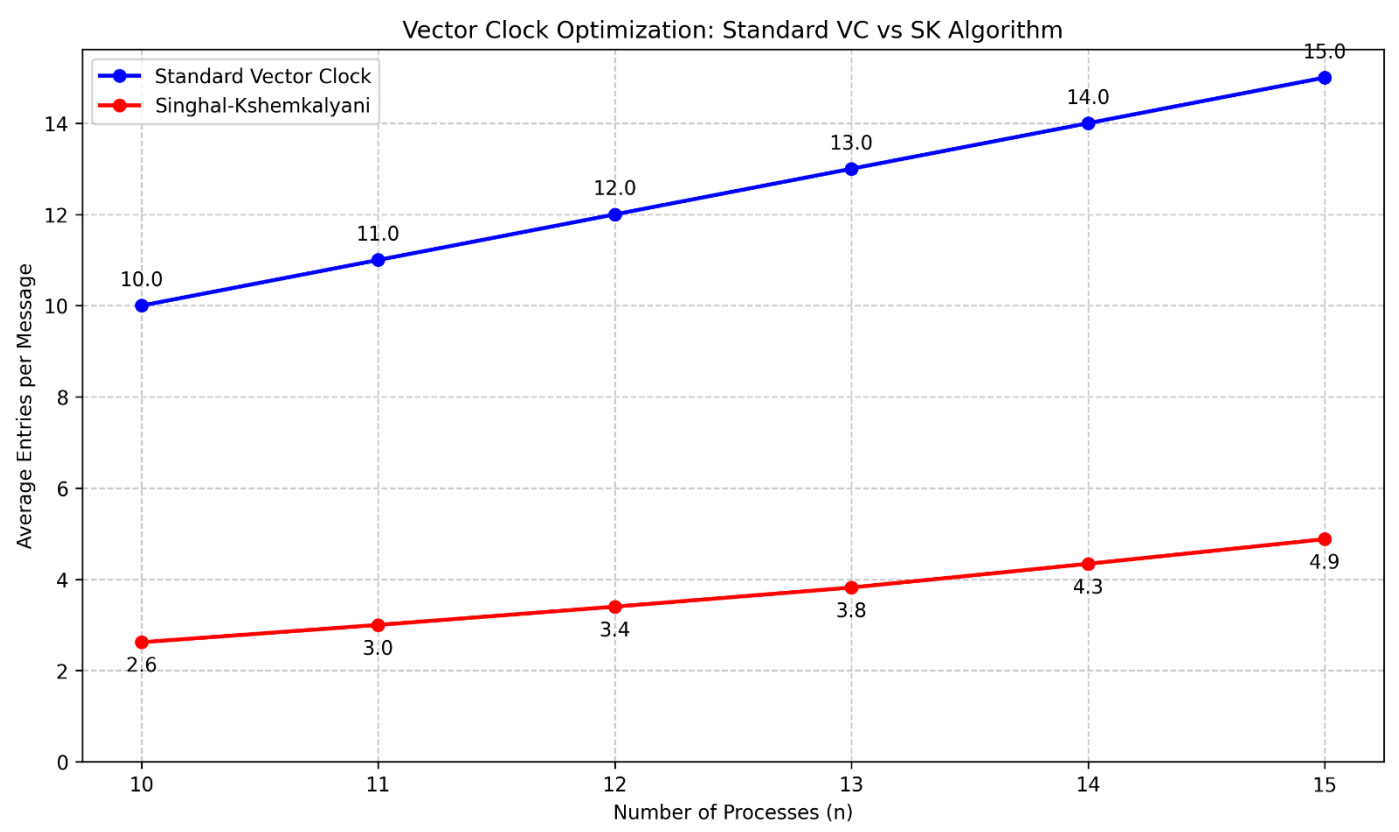
If  $n=3$ , process 1's neighbors are [2,3], process 2's neighbors are [1,3], and process 3's neighbors are [1,2]

If  $n=4$ , process 1's neighbors are [2,3,4], process 2's neighbors are [1,3,4], etc.

### 3. Results

Number of Processes (n)	Standard VC (entries/msg)	SK Optimization (entries/msg)	Reduction (%)
10	10.0	2.6	74.0%
11	11.0	3.0	72.7%
12	12.0	3.4	71.7%
13	13.0	3.8	70.8%
14	14.0	4.3	69.3%
15	15.0	4.9	67.3%

Graph:



Observations:

Standard Vector Clock Behaviour:

- Shows perfectly linear growth with  $O(n)$  complexity
- Average entries per message equals  $n$  (as expected)
- Predictable but increasingly expensive as system scales

### Singhal-Kshemkalyani Optimization:

- Demonstrates sub-linear growth
- Average entries range from 2.6 to 4.9 (approximately 26-33% of  $n$ )
- Maintains strong consistency while significantly reducing overhead

### Scalability Analysis:

- The gap between algorithms widens as  $n$  increases
- At  $n=10$ : SK saves 7.4 entries per message
- At  $n=15$ : SK saves 10.1 entries per message
- **Savings increase by ~36% as system scales from 10 to 15 processes**

## 4. Factors affecting performance

### Positive factors (reduce entries):

- Sparse communication patterns
- Lower  $\alpha$  values (fewer internal events means less frequent local clock updates)
- Localized communication (processes primarily communicate with neighbours)

### Negative factors (increase entries):

- Dense communication graphs
- High  $\alpha$  values
- Frequent internal events causing many local updates

For a system with  $n$  processes:

- **Best case SK:**  $O(1)$  entries when only local clock changes
- **Worst case SK:**  $O(n)$  entries when all clocks have been updated
- **Average case SK:**  $O(k)$  where  $k$  is the average number of processes in causal dependency chains

## 5. Conclusion

- **Significant Overhead Reduction:** SK optimization reduces message overhead by 67-74% across all tested configurations
- **Scalability:** The optimization becomes increasingly valuable as system size grows